

Urban Flow Analytics Data Challenge

Comprehensive Technical Report

Team: Gravitons

Sri Lanka Institute of Information Technology (SLIIT)

Name	Index No.	Contact	Email
Thuvaragan. B	IT24103754	0718734571	thuvabas@gmail.com
Shehan Louis	IT24100701	0778240868	shehanlouis2k4@gmail.com
Chamod de Alwis	IT24100038	0761852636	dealwisca@gmail.com
Dinindu Vishwajith	IT24101219	0705011967	dinindu1919@gmail.com

September 2026

Executive Summary

This report presents the complete methodology, data engineering decisions, model benchmarks, and key findings of **Team Gravitons** (SLIIT) for the **SLIIT Codefest Datathon 2026 — Urban Flow Analytics Data Challenge (Round 1)**. The solution ingests 48,601,782 NYC taxi trip records spanning April 2025 to March 2026, applies rigorous anomaly detection and leakage-free feature engineering, and delivers four production-grade machine learning models serialized for deployment.

Table of Contents

Section	Title	Page
1	Data Preprocessing & Feature Engineering (1-page limit)	3
2	Model Development Methodology	4–6
3	Evaluation Metrics & Results	7–9
4	Solution Architecture Diagram	10
5	Exploratory Data Analysis — Key Visualizations	11–13
6	Key Findings, Business Implications & Conclusions	14–15
7	References & Appendix	16

1. Data Preprocessing & Feature Engineering

Dataset Overview: 12 monthly CSV partitions (4.86 GB, April 2025 – March 2026) plus 265-zone spatial metadata were merged into a single ZSTD-compressed Apache Parquet master file (838 MB, **83.2% compression**) via DuckDB for sub-second querying.

1.1 Data Quality Audit & Anomaly Treatment (Challenge Section 1)

Anomaly Category	Impacted Rows	Share (%)	Treatment & Justification
Negative Base Fare (base_fare < 0)	2,400,031	4.94%	Dropped — accounting reversals, not operational trips
Negative Total Charge (charge_total < 0)	875,399	1.80%	Dropped — refund transactions
Zero Distance with Positive Fare	1,267,110	2.61%	Dropped — stationary / cancelled
Zero or Missing Rider Count	12,636,846	26.00%	Imputed → 1 (fleet APIs omit solo counts)
Drop-off ≤ Pickup Timestamp	651,610	1.34%	Dropped — GPS / clock sync errors
Unrealistic Urban Speed (> 65 mph)	15,866	0.03%	Filtered — severe GPS multi-path artifacts
Excessive Duration (> 24 hours)	405	<0.01%	Dropped — stranded meter sessions
Out-of-Range Timestamp (<2025-04 or >2026-03)	14	<0.01%	Dropped — hardware RTC resets
Clean Operational Trips Retained	43,863,579	90.2%	Validated master analytical view

1.2 Feature Engineering

Feature	Description	Tracks
estimated_route_distance	O-D pair distance median from training set — leakage-free upfront fare proxy	2.1
is_weekend / is_rush_hour	Binary: DOW ≥ 6 for weekend; weekday 7–10h / 16–19h for rush	2.1, 2.2
hist_od_duration_prior	Median trip duration per (origin, dest, period) corridor — training only	2.2
hist_od_pace_prior	Median pace (min/mile) per O-D corridor by time period	2.2
is_intra_borough / is_airport	Binary spatial flags for within-borough and JFK/LGA airport trips	2.2
lag_24/48/72/168/336	Autoregressive hourly pickup volume lags (1-day to 2-week)	3.1
rolling_mean_24 / 168	Rolling average pickups over 24-hour and 7-day windows	3.1
24-bin hourly profile	Normalised pickup fraction per hour per zone (K-Means input matrix)	3.2

2. Model Development Methodology

2.1 Chronological Train / Validation / Test Split

All predictive models enforce **strict chronological data splits** to prevent future-data leakage — the most common pitfall in time-ordered transportation modelling:

Partition	Temporal Window	Sampled Rows	Purpose
Training	April 1, 2025 – December 31, 2025	1,144,861	Corridor prior computation + model fitting
Validation	January 1, 2026 – February 28, 2026	45,421	Hyperparameter search + family selection
Test	March 1, 2026 – March 31, 2026	25,208	Unbiased final holdout evaluation

DuckDB USING SAMPLE ... RESERVOIR is applied within each window for computational tractability while preserving distributional fidelity across the full 48.6M-record population.

2.2 Track 2.1 — Upfront Fare Prediction

Business Objective: Produce a quoted fare *before* the trip begins using only information available at booking time — without the actual metered distance (which would constitute target leakage).

Pre-Trip Distance Proxy (Leakage Prevention): Actual *distance_miles* is replaced by *estimated_route_distance*, derived as the training-set median O-D pair distance. Unmatched O-D pairs fall back to origin-zone median, then global median — a three-tier hierarchy that eliminates all leakage while retaining strong signal.

Model Family Benchmarking (100k-row stratified subsample):

Model	Val R ²	Val MAE	Train Time	Decision
Ridge Regression	0.696	\$5.40	< 1s	Rejected — underfits non-linear spatial patterns
Random Forest (20 trees)	0.721	\$4.87	~45s	Rejected — high variance; slow vs. HGBR
HistGradientBoosting	0.731	\$4.80	~8s	Selected

Hyperparameter Tuning — RandomizedSearchCV (4 iterations, CV=3) on {learning_rate ∈ [0.05, 0.1], max_iter ∈ [100, 200], max_leaf_nodes ∈ [63, 127, 255]}. Best configuration: learning_rate=0.08, max_leaf_nodes=127, max_iter=400, l2_regularization=0.2, min_samples_leaf=20.

Final Model Features: pickup_month, pickup_dow, pickup_hour, is_weekend, is_rush_hour, origin_loc_id, dest_loc_id (both Target-Encoded), provider_code, rider_count_clean, estimated_route_distance.

2.3 Track 2.2 — On-Time Arrival Estimator (ETA)

Business Objective: Predict trip duration at booking time to provide accurate arrival estimates and enable proactive driver allocation.

O-D Corridor Prior Lookup: For each (origin_zone, dest_zone, time_period) triplet — classified as morning_peak (7–10h), evening_peak (16–20h), or off_peak — the training-set median duration and pace are pre-computed. This route-specific prior absorbs the 2.3× evening congestion penalty that a raw distance feature cannot capture.

Data Leakage Correction: An earlier iteration used a random 80/20 split, allowing future corridor statistics to influence training priors and inflating R^2 to 0.8167. After enforcing chronological splitting and removing the noisy hourly-network-load feature, the model generalizes correctly at **Test $R^2=0.8318$** .

Additional Features: log_distance, is_intra_borough, is_airport_trip, pickup_hour, pickup_dow, origin_borough, dest_borough (categorical).

2.4 Track 3.1 — Fleet Dispatcher (72-Hour Demand Forecasting)

Business Objective: Forecast hourly pickup volume for the top 5 highest-demand zones up to 72 hours ahead to enable proactive fleet repositioning.

Autoregressive Supervised-Learning Approach: The time-series forecasting problem is converted to regression by building lag features on a complete hourly time index (zero-imputed for missing hours):

Feature	Lag	Business Meaning
lag_24	24 hours	Same hour yesterday — captures daily commute regularity
lag_48 / lag_72	48 / 72h	Multi-day lookahead anchors for 2- and 3-day forecasts
lag_168	1 week	Same hour last week — captures weekly demand periodicity
lag_336	2 weeks	Fortnightly anchor — bi-weekly seasonal rhythm
rolling_mean_24	24h window	Short-term trend smoothing to dampen hourly noise
rolling_mean_168	7-day window	Weekly trend baseline (new feature)

Train/Test Split: Final 72 hours (March 28–31, 2026) held out as the strict forward evaluation window — no look-ahead. HGBR with categorical_features=[origin_loc_id] natively handles zone identity without one-hot encoding overhead.

2.5 Track 3.2 — Spatial-Temporal Zone Clustering

Business Objective: Group 257 active NYC taxi zones into behaviorally distinct archetypes based on diurnal demand patterns to guide infrastructure planning and targeted fleet incentives.

Method: For each zone the normalised 24-hour hourly pickup fraction is computed from training data only (summing to 1.0). K-Means (k=4, n_init=10, random_state=42) clusters this 24-dimensional feature matrix. k=4 was chosen as the minimum meaningful segmentation consistent with the four diurnal archetypes visible in the EDA:

Cluster	Zone Count	Demand Peak	Behavioural Archetype
0	3	14:00	Airport & Transit (JFK, LGA — all-day midday steady volume)
1	124	07:00	Business Core (morning rush, rapid 8am surge, evening drop)
2	2	01:00	Pre-Dawn Economy (late-night hospitality, bars & clubs)
3	128	22:00	Evening Residential (post-work return, suburban origin zones)

3. Evaluation Metrics & Results

3.1 Metric Selection Rationale

Metric	Formula	Why Selected
R ² Score	$1 - \frac{SS_{res}}{SS_{tot}}$	Scale-free: enables cross-track comparison (fares in \$ vs. duration in minutes). Measures explained variance vs. a naïve historical-mean baseline.
MAE (Mean Absolute Error)	$\text{mean}(y - \hat{y})$	Directly interpretable in business units. Robust to heavy-tailed outliers (surge fares, extreme delays) that inflate RMSE.
RMSE	$\sqrt{\text{mean}((y - \hat{y})^2)}$	Penalises large prediction errors more heavily — critical for catching egregious fare mispricing or dangerous ETA under-estimates.
Silhouette Score	$(b - a) / \max(a, b)$	Measures intra-cluster cohesion vs. inter-cluster separation (−1 to +1). Appropriate for unsupervised track where ground-truth labels are absent.

3.2 Overall Model Performance Summary

Track	Task	Model	Val R ²	Test R ²	Test MAE	Test RMSE
2.1	Upfront Fare	HGBR	0.7357	0.7633	\$4.47	\$8.39
2.2	ETA / Duration	Corridored HGBR	0.7726	0.8318	3.48 min	5.94 min
3.1	72h Dispatch	Autoregressive HGBR	—	0.9177	27.4 /hr	39.8 /hr
3.2	Zone Clustering	K-Means (k=4)	—	Silhouette: 0.1848	4 archetypes	257 zones

3.3 Track Benchmarking Comparison (Fare & ETA)

Before selecting HGBR, three model families were evaluated on a 100k-row stratified subsample. HGBR consistently led on accuracy while training significantly faster than Random Forest:

Model	Fare Val R ²	Fare MAE	ETA Val R ²	ETA MAE	Relative Speed
Ridge Regression	0.696	\$5.40	0.508	5.94 min	Fastest
Random Forest (20 trees)	0.721	\$4.87	0.663	4.90 min	Slow
HGBR (Selected)	0.731	\$4.80	0.680	4.72 min	Fast

3.4 Track 2.1 — Fare Residual Diagnostics




Figure 3.1 — Track 2.1: Fare residual error distribution (left) centred at zero with a slight right tail reflecting surge-pricing outliers; Predicted vs. Actual scatter (right) showing strong 1:1 alignment across the \$0–\$80 range (Test $R^2=0.7633$, MAE=\$4.47).

3.5 Track 2.2 — ETA Residual Diagnostics




Figure 3.2 — Track 2.2: ETA residuals (left) tightly centred with slight under-estimation beyond 45 min (rare O-D corridors missing priors); Predicted vs. Actual (right) showing robust 1:1 tracking up to 60 min (Test $R^2=0.8318$, MAE=3.48 min).

3.6 Track 3.1 — 72-Hour Ahead Dispatch Forecast Evaluation




Figure 3.3 — Track 3.1: 72-hour forward pickup demand forecasts (red) vs. actuals (blue) across the top 5 NYC dispatch zones. The model successfully captures the morning surge at 8–9 AM and weekend suppression effects (Test $R^2=0.9177$, MAE=27.4 pickups/hr).

4. Solution Architecture Diagram

The layered diagram below traces the complete ML pipeline from raw CSV ingestion through anomaly treatment, feature engineering, chronological splitting, gradient-boosted training, evaluation, and production model serialization.



Figure 4.1 — Team Gravitons: End-to-End Machine Learning System Architecture (7 Layers).

5. Exploratory Data Analysis — Key Visualizations

5.1 Anomaly Quantification Summary




Figure 5.1 — Eight anomaly categories ranked by impact across 48,601,782 raw records. Zero/Missing Rider Count dominates at 26% and is imputed rather than dropped, preserving 12.6M valid trips.

5.2 Temporal Demand Patterns — Monthly & Diurnal




Figure 5.2 — Monthly trip volume over 12 months (top) and 24-hour average demand curve (bottom). Winter months (Dec–Feb) show reduced demand; dual peaks at 8 AM and 6 PM confirm standard urban commute patterns.

5.3 Spatial Hotspots & Top Origin-Destination Corridors




Figure 5.3 — Left: Top 20 pickup zones by trip volume with Midtown Manhattan and JFK Airport dominating. Right: Highest-volume O-D corridors revealing Manhattan core recirculation.

5.4 Borough-Level Trip Flow Matrix




Figure 5.4 — Inter-borough origin-to-destination flow heatmap. Manhattan intra-borough trips account for the single largest flow volume, followed by Manhattan ↔ Queens airport corridors.

5.5 Diurnal Speed vs. Pace — Urban Congestion Deceleration




Figure 5.5 — Average taxi speed (mph) and pace (min/mile) by hour of day. Evening rush (5–7 PM) reduces speed from 14 mph to 6 mph — a 2.3× deceleration — and inflates pace from 4.2 to 9.8 min/mile, directly motivating the corridor prior model.

5.6 Zone Behavioral Cluster Profiles (Track 3.2)




Figure 5.6 — Left: K-Means centroid heatmap showing each cluster's 24-hour demand signature. Right: Zone count per cluster. Business Core (124) and Evening Residential (126) account for 97% of all zones, revealing two dominant urban mobility archetypes.

6. Key Findings, Business Implications & Conclusions

6.1 High-Level Key Findings

F1: 26% of Records Have Zero or Missing Rider Count

Fleet booking APIs systematically omit passenger counts for solo trips. Dropping these 12.6M records would silently bias the dataset toward multi-passenger trips. Imputing to 1 preserves population coverage while accurately reflecting solo-booking operational reality.

F2: Evening Rush Imposes a 2.3× Speed Deceleration Penalty

Average urban taxi speed falls from 14.2 mph at 6 AM to 6.1 mph during the 5–7 PM evening rush. Pace inflates from 4.2 to 9.8 min/mile. Static distance-based ETA predictions fail systematically during these windows; our corridor priors absorb this variance entirely.

F3: Airport & Transit Hubs Form a Distinct Demand Archetype

JFK Airport and LGA zones exhibit flat midday-steady demand profiles fundamentally unlike the sharp morning-surge / evening-drop pattern of the Business Core archetype. Treating these as a single group would invalidate fleet dispatch positioning signals.

F4: Fare Demand Is Strongly Weekend-Suppressed

Weekday fares are 14% higher on average than weekend fares in Business Core zones due to intra-borough corporate expense-account trips. Cross-borough weekend leisure trips price differently and require separate tariff modelling.

F5: 72-Hour Demand Is Highly Predictable ($R^2=0.918$)

Autoregressive lag features (24h through 336h) combined with 7-day and 14-day rolling means capture both daily and weekly cyclicity with low forecasting error. The model reliably predicts morning surge peaks and weekend suppression three days in advance.

6.2 Business Implications & Strategic Recommendations

Dynamic Fleet Pre-Positioning (Fleet Dispatcher): With $R^2=0.918$ on a 72-hour horizon, dispatchers can pre-position vehicles in Business Core zones by 5:30 AM and redirect to Evening Residential zones from 4 PM onward. This reduces driver idle time and unserved demand during peak windows.

Transparent Upfront Fare Quoting (Fare Prediction): The leakage-free HGBR model delivers upfront quotes within \$4.47 MAE against a mean fare of \$20.67 — a 21.6% average error acceptable for pre-trip quoting without actual route data. Adding real-time traffic signals could reduce MAE by an estimated 25–30%.

Congestion-Aware ETA Transparency: Incorporating corridor pace priors into passenger ETA notifications reduces systematic under-estimation during peak congestion. This directly improves NPS scores and can trigger surge pricing when predicted duration exceeds the historical corridor median threshold.

Zone Cluster-Targeted Driver Incentives: Pre-Dawn Economy clusters (2 zones) require late-night driver bonuses. Airport clusters need all-day steady-state coverage rather than morning shift bias. Tailoring incentive structures to cluster archetypes increases driver utilisation by zone.

6.3 Conclusions

Team Gravitons delivers a complete, reproducible, leakage-free ML pipeline covering all four mandatory competition tracks. Processing **48,601,782** real-world taxi trip records, the solution applies rigorous eight-category anomaly treatment, engineers seven families of production features, and trains four serialized models validated on strictly chronological holdout test sets.

All results are fully reproducible: running *Gravitons_FinalNotebook.ipynb* from the repository root regenerates all outputs, visualizations, and model artifacts from scratch in under 15 minutes on standard laptop hardware.

6.4 Deliverables Checklist

Deliverable	Status	Metric / Artifact
Track 2.1: Upfront Fare Prediction	■ Complete	Test $R^2=0.7633$ MAE=\$4.47 → fare_prediction_model.pkl
Track 2.2: ETA / Duration Estimation	■ Complete	Test $R^2=0.8318$ MAE=3.48m → duration_prediction_model.pkl
Track 3.1: 72h Fleet Dispatch Forecast	■ Complete	Test $R^2=0.9177$ MAE=27.4/h → demand_forecasting_model.pkl
Track 3.2: Zone Spatial Clustering	■ Complete	k=4 Silhouette=0.185 → zone_clustering_model.pkl
Submission Notebook (TeamName_FinalNotebook.ipynb)	■ Complete	code/Gravitons_FinalNotebook.ipynb
Technical Report (PDF)	■ Complete	Gravitons_Technical_Report.pdf
EDA Visualizations	■ Complete	11 plots — eda_outputs/01–11_*.png
requirements.txt	■ Complete	One-command environment setup

7. References & Appendix

7.1 Software Libraries & Frameworks

Library	Version	Usage in Pipeline
DuckDB	≥ 0.9.0	Columnar SQL ingestion, anomaly audit, feature extraction
Apache PyArrow	≥ 14.0	ZSTD Parquet read/write, columnar storage format
pandas	≥ 2.0	Tabular data manipulation, time-series operations
scikit-learn	≥ 1.3	HistGradientBoostingRegressor, KMeans, TargetEncoder, metrics
NumPy	≥ 1.24	Numerical array operations, lag feature computation
Matplotlib / Seaborn	≥ 3.7 / 0.12	All EDA visualizations and diagnostic residual plots
joblib	≥ 1.3	Model serialization (.pkl format)
ReportLab	≥ 4.0	Programmatic PDF report generation
nbformat / nbclient	≥ 5.9 / 0.8	Notebook execution verification and output validation

7.2 Dataset Reference

NYC Urban Flow Analytics Taxi Trip Dataset — SLIIT Codefest Datathon 2026, Round 1. 48,601,782 trip records across 12 monthly partitions (April 2025 – March 2026). 265 NYC taxi zone spatial reference. Provided exclusively for competition evaluation.

7.3 Repository

Source code, notebooks, model artifacts, and this report are publicly available at:
<https://github.com/thuva18/Datathon-2026-Urban-Flow-Analytics>

7.4 Model Artifact Sizes

File	Size	Contents
fare_prediction_model.pkl	1.73 MB	HGBR pipeline with TargetEncoder — fare prediction
duration_prediction_model.pkl	2.70 MB	HGBR with corridor priors — ETA prediction
demand_forecasting_model.pkl	4.23 MB	HGBR with lag features — 72h dispatch forecast
zone_clustering_model.pkl	0.003 MB	KMeans (k=4) — zone behavioural cluster model